

Alex Candidate - Applied AI Engineer

Experienced machine learning engineer specializing in LLM deployment and RAG architectures.

- Deployed vLLM inference server on Kubernetes with 99.9% uptime.